

Propensity Scores

Lucy D'Agostino McGowan

Wake Forest University


Observational Studies









Confounding




Confounding




What part of the DAG do we want to try to deal with?



What part of the DAG do we want to try to deal with?



What part of the DAG do we want to try to deal with?



Which relationship should we model?

- Paths into the exposure: propensity scores
- Paths into the outcome: outcome models
- Both: doubly robust estimators

The hidden experiment

- Imagine the exposure had been randomized with a probability of 0.5
- Then the arrows into the exposure wouldn't exist
- In an experiment, the propensity score is *known*:
`rbinom(1, 1, 0.5)`

The hidden experiment

- Now imagine randomizing the exposure *within levels of the confounders*
- Still a randomized experiment, but the probability of exposure depends on the covariates
- This gives us conditional exchangeability

The hidden experiment

- In observational data, these probabilities are not known
- But *something* determined who was exposed: what was that decision process?
- Domain knowledge and statistics can estimate it: the **propensity score**

Propensity scores

Rosenbaum and Rubin showed that in observational studies, conditioning on **propensity scores** can lead to unbiased estimates of the exposure effect

- 1 There are no unmeasured confounders
- 2 Every subject has a nonzero probability of receiving either exposure

Propensity scores

- Fit a **logistic regression** predicting exposure using known covariates
- Each individual's predicted value is their **propensity score**

Propensity scores

```
1 library(tidyverse)
2 library(broom)
3 library(causalworkshop)
4 library(halfmoon)
```



Propensity scores

```
1 glm(  
2   exposure ~ confounder_1 + confounder_2 + confounder_3 + ...,  
3   data = df,  
4   family = binomial()  
5 )
```

Propensity scores

```
1 glm(  
2   exposure ~ confounder_1 + confounder_2 + confounder_3 + ...,  
3   data = df,  
4   family = binomial()  
5 ) |>  
6   augment(type.predict = "response", data = df)
```

Propensity scores



Propensity scores

```
1 propensity_model <- glm(  
2   net ~ income + health + temperature,  
3   data = net_data,  
4   family = binomial()  
5 )  
6  
7 net_data_ps <- propensity_model |>  
8   augment(type.predict = "response", data = net_data)
```


Calibration



If we look at all the households with a predicted probability of 0.3, about 30% of them should use a net

```
1 plot_model_calibration(  
2     propensity_model,  
3     method = "windowed",  
4     window_size = 0.15  
5 )
```

Calibration




ROC curve



How well does the model discriminate between the exposed and the unexposed?

```
1 net_data_ps |>  
2   check_model_roc_curve(net, .fitted) |>  
3   plot_model_roc_curve()
```

ROC curve



AUC

```
1 check_model_auc(net_data_ps, .exposure = net, .fitted = .fitted)
```

```
# A tibble: 1 × 2  
  method      auc  
  <chr>      <dbl>  
1 observed 0.600
```

- Perfect discrimination would be a violation of positivity
- Values between 0.6 and 0.8 usually mean we've picked up enough signal to adjust for

AUC

- There is no clear cutoff for a “good” AUC in causal inference
- A randomized trial has an expected AUC of 0.5
- We are not trying to optimize the AUC; we use it as a diagnostic

Example: The Seven Dwarfs Mine Train




Photo by Anna [CC-BY-SA-4.0](#)

Historically, guests who stayed in a Walt Disney World resort hotel were able to access the park during “Extra Magic Hours” during which the park was closed to all other guests.

These extra hours could be in the morning or evening.

The Seven Dwarfs Mine Train is a ride at Walt Disney World’s Magic Kingdom. Typically, each day Magic Kingdom may or may not be selected to have these “Extra Magic Hours”.

We are interested in examining the relationship between whether there were “Extra Magic Hours” in the morning and the average posted wait time for the Seven Dwarfs Mine Train the same day between 9 AM and 10 AM in 2018.



Your Turn

Using the confounders identified in the previous DAG, fit a propensity score model for `park_extra_magic_morning`

Check the calibration, ROC curve, and AUC of your model

Stretch: Create two histograms, one of the propensity scores for days with Extra Magic Hours in the morning and one for those without